

```
In [57]: import pandas as pd
import numpy as np
import seaborn as sns
from sklearn .model_selection import train_test_split
from matplotlib import pyplot as plt
```

```
In [58]: df= pd.read_csv(r"C:\Users\teppa\Downloads\Advertising.csv")
df
```


Out[58]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9
...
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

200 rows × 4 columns

```
In [59]: plt.scatter(df['TV'],df['Sales'])
```

Out[59]: <matplotlib.collections.PathCollection at 0x201e070aaa0>



```
In [60]: x=df[['TV']]
y=np.array(df['Sales']).reshape(-1,1)
```

```
In [61]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
x_train
```


Out[61]:

	TV
195	38.2
29	70.6
88	88.3
14	204.1
79	116.0
...	...
44	25.1
140	73.4
149	44.7
73	129.4
45	175.1

140 rows × 1 columns

```
In [62]: sns.pairplot(df,x_vars=["TV","Radio","Newspaper"],y_vars='Sales',height=5,aspect=0.8,kind='r')
```

Out[62]: <seaborn.axisgrid.PairGrid at 0x201da7a46d0>



```
In [63]: len(x_train)
```

Out[63]: 140

```
In [64]: len(x_test)
```

Out[64]: 60

```
In [65]: len(y_train)
```

```
Out[65]: 140
```

```
In [66]: len(y_test)
```

```
Out[66]: 60
```

```
In [67]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()
```

```
In [68]: lr.fit(x_train,y_train)
```


```
Out[68]: ▾ LinearRegression  
LinearRegression()
```

```
In [69]: lr.predict(x_test)
```

```
Out[69]: array([[13.03303058],
 [19.52785079],
 [ 9.94730588],
 [20.16363066],
 [17.84522648],
 [ 9.66230111],
 [ 7.73851892],
 [22.59713291],
 [10.79683932],
 [ 7.23976057],
 [17.30810211],
 [18.25081019],
 [20.34449907],
 [11.20242303],
 [22.55876688],
 [ 9.17998535],
 [ 7.99611938],
 [23.06848695],
 [18.50292979],
 [14.23333913],
 [18.99072642],
 [16.24481509],
 [13.6194827 ],
 [20.92547032],
 [17.84522648],
 [ 9.03748296],
 [21.64346311],
 [12.64388945],
 [11.30107853],
 [21.55028847],
 [20.17459238],
 [20.15814979],
 [16.64491794],
 [12.35888469],
 [11.79435601],
 [13.64688701],
 [ 9.45402839],
 [ 8.03996627],
 [11.15309528],
 [17.6314729 ],
 [15.96529118],
 [21.41326695],
 [10.68174124],
 [13.43861429],
 [17.96580542],
 [18.56321926],
 [22.01616165],
 [14.6444037 ],
 [ 7.95775336],
 [13.08235833],
 [18.71668337],
 [14.66084628],
 [19.7361235 ],
 [18.87562834],
 [ 7.94131077],
 [19.62650628],
 [12.90148992],
 [ 7.31101177],
 [18.89755178],
 [11.19694217]])
```

```
In [70]: print("Rgression:",lr.score(x_test,y_test))
y_pred=lr.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```

Rgression: 0.8191900194075661



```
In [71]: lr.score(x_test,y_test)
```

Out[71]: 0.8191900194075661

```
In [72]: df100=df[:][:180]
```

```
In [73]: x=df100[["TV"]]
y=df100[["Sales"]]
```

```
In [74]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
x_train
```

Out[74]:

	TV
100	222.4
132	8.4
70	199.1
3	151.5
144	96.2
...	...
148	38.0
142	220.5
15	195.4
66	31.5
74	213.4

126 rows × 1 columns

```
In [75]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
```

```
In [76]: feature=["TV"]
target=["Sales"]
```

```
In [77]: lr.fit(x_train,y_train)
```

Out[77]:

▼ LinearRegression

LinearRegression()

```
In [78]: lr.predict(x_test)
```


Out[78]: array([23.08838931, 16.55080001, 21.81657159, 17.43772552, 11.29060656,
19.02191951, 9.41077073, 23.14417079, 15.14510674, 18.83784063,
8.32861005, 10.86666733, 18.46410472, 7.22971491, 14.24144679,
17.40983478, 21.72174308, 15.28456044, 19.18926395, 7.41379379,
12.16637578, 20.49455055, 20.17101797, 7.87120192, 7.40821565,
11.9711406 , 20.21564316, 12.3671891 , 19.42354616, 11.79263987,
8.25609412, 14.75463639, 13.05330129, 22.586356 , 7.15719899,
8.35650078, 9.91280404, 9.42192703, 10.25307106, 18.62029287,
11.93209357, 14.14661827, 17.24249034, 18.22424437, 9.33267666,
18.94382544, 14.89966824, 14.27491567, 18.02343104, 7.88793636,
22.78716932, 11.11768398, 21.17508459, 12.65725279])

```
In [79]: lr.score(x_test,y_test)
```

Out[79]: 0.7555145093619565

```
In [80]: print("Rgression:",lr.score(x_test,y_test))
y_pred=lr.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```

Rgression: 0.7555145093619565



```
In [81]: from sklearn.linear_model import Ridge,RidgeCV,Lasso
from sklearn.preprocessing import StandardScaler
```

```
In [82]: ridge=Ridge(alpha=10)
ridge.fit(x_train,y_train)
```

```
Out[82]:
```


▼	Ridge
	Ridge(alpha=10)

```
In [83]: train_score_ridge=ridge.score(x_train,y_train)
test_score_ridge=ridge.score(x_test,y_test)
print("\nRidge Model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.8349938083289384
The test score for ridge model is 0.7555153346434035

```
In [84]: plt.figure(figsize=(10,10))
plt.plot(feature,ridge.coef_,alpha=0.7,linestyle='None',marker='*',markersize=5,color='red',)
plt.plot(feature,lr.coef_,alpha=0.4,linestyle="none",marker='o',markersize=7,color='green',)
plt.xticks(rotation=90)
plt.legend()
plt.show()
```





```
In [85]: print("\nLasso Model:\n")
lasso=Lasso(alpha=10)
lasso.fit(x_train,y_train)
train_score_ls=lasso.score(x_train,y_train)
test_score_ls=lasso.score(x_test,y_test)
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Lasso Model:

The train score for ridge model is 0.8349938083289384
The test score for ridge model is 0.7555153346434035

```
In [86]: pd.Series(lasso.coef_,feature).sort_values(ascending=True).plot(kind="bar")
```

Out[86]: <Axes: >




```
In [87]: from sklearn.linear_model import LassoCV
```

```
In [88]: lasso_cv=LassoCV(alphas=[0.0001,0.001,0.01,1,10],random_state=0).fit(x_train,y_train)
```

```
In [89]: print(lasso_cv.score(x_train,y_train))
print(lasso_cv.score(x_test,y_test))
```

0.8349938084374653
0.7555145280061444

```
In [90]: plt.figure(figsize=(10,10))
plt.plot(feature,ridge.coef_,alpha=0.7,linestyle='None',marker='*',markersize=5,color='red',label='Ridge')
plt.plot(feature,lasso_cv.coef_,alpha=0.4,linestyle='none',marker='d',markersize=6,color='blue',label='Lasso')
plt.plot(feature,lr.coef_,alpha=0.4,linestyle="none",marker='o',markersize=7,color='green',label='Linear Regression')
plt.xticks(rotation=90)
plt.legend()
plt.title("Comparsion plot of Ridge, lasso and Linear regression model")
plt.show()
```



```
In [91]: from sklearn.linear_model import RidgeCV
```

```
In [92]: ridge_cv=RidgeCV(alphas=[0.0001,0.001,0.01,1,10]).fit(x_train,y_train)
```

```
In [93]: print("The train score for ridge for ridge model is {}".format(ridge_cv.score(x_train,y_train))
print("The test score for ridge for ridge model is {}".format(ridge_cv.score(x_test,y_test)))
```

The train score for ridge for ridge model is 0.8349938083289381
The test score for ridge for ridge model is 0.7555153346448448

```
In [94]: from sklearn.linear_model import ElasticNet
```

```
In [95]: regr=ElasticNet()
regr.fit(x,y)
print(regr.coef_)
print(regr.intercept_)
```

[0.05491733]
7.065459085981752

```
In [102]: y_predt_elastic=regr.predict(x_train)
```

```
In [104]: mean_squared_error=np.mean((y_predt_elastic-y_train)**2)
print("Mean squared Error on test set",mean_squared_error)
```

Mean squared Error on test set 4.284671287369669

```
In [ ]:
```

```
In [ ]:
```