

In [76]:

```
1 import numpy as np
2 import pandas as pd
3 import seaborn as sns
4 import matplotlib.pyplot as plt
5 from sklearn import preprocessing, svm
6 from sklearn.model_selection import train_test_split
7 from sklearn.linear_model import LinearRegression
```

```
In [77]: 1 df=pd.read_csv(r"C:\Users\teppa\Desktop\AHISAI\bottle.csv")
        2 df
```

C:\Users\teppa\AppData\Local\Temp\ipykernel\_1224\3415337971.py:1: DtypeWarning: Columns (47,73) have mixed types. Specify dtype option on import or set low\_memory=False.  
df=pd.read\_csv(r"C:\Users\teppa\Desktop\AHISAI\bottle.csv")

Out[77]:

	Cst_Cnt	Btl_Cnt	Sta_ID	Depth_ID	Depthm	T_degC	Salnty	O2ml_L	STheta	O2Sat	...	R_PHAEO	R_PRES	R_
0	1	1	054.0 056.0	19-4903CR-HY-060-0930-05400560-0000A-3	0	10,500	33,4400	NaN	25,64900	NaN	...	NaN	0	
1	1	2	054.0 056.0	19-4903CR-HY-060-0930-05400560-0008A-3	8	10,460	33,4400	NaN	25,65600	NaN	...	NaN	8	
2	1	3	054.0 056.0	19-4903CR-HY-060-0930-05400560-0010A-7	10	10,460	33,4370	NaN	25,65400	NaN	...	NaN	10	
3	1	4	054.0 056.0	19-4903CR-HY-060-0930-05400560-0019A-3	19	10,450	33,4200	NaN	25,64300	NaN	...	NaN	19	
4	1	5	054.0 056.0	19-4903CR-HY-060-0930-05400560-0020A-7	20	10,450	33,4210	NaN	25,64300	NaN	...	NaN	20	
...	...	...	...	...	...	...	...	...	...	...	...	...	...	
864858	34404	864859	093.4 026.4	20-1611SR-MX-310-2239-09340264-0000A-7	0	18,744	33,4083	5,805	23,87055	108,74	...	0,18	0	
864859	34404	864860	093.4 026.4	20-1611SR-MX-310-2239-09340264-0002A-3	2	18,744	33,4083	5,805	23,87072	108,74	...	0,18	2	
864860	34404	864861	093.4 026.4	20-1611SR-MX-310-2239-09340264-0005A-3	5	18,692	33,4150	5,796	23,88911	108,46	...	0,18	5	
864861	34404	864862	093.4 026.4	20-1611SR-MX-310-2239-09340264-0010A-3	10	18,161	33,4062	5,816	24,01426	107,74	...	0,31	10	
864862	34404	864863	093.4 026.4	20-1611SR-MX-310-2239-09340264-0015A-3	15	17,533	33,3880	5,774	24,15297	105,66	...	0,61	15	

864863 rows × 74 columns



```
In [78]: 1 df=df[['Salnty','T_degC']]
         2 df.columns=['sal','tem']
```


```
In [79]: 1 df.head()
```

```
Out[79]:
```

	sal	tem
0	33.440	10.50
1	33.440	10.46
2	33.437	10.46
3	33.420	10.45
4	33.421	10.45

```
In [80]: 1 sns.lmplot(x='sal',y='tem',data=df,order=2,ci=None)
```

```
Out[80]: <seaborn.axisgrid.FacetGrid at 0x17c29516aa0>
```



```
In [81]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 864863 entries, 0 to 864862
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype  
---  -
0    sal      817509 non-null    float64
1    tem      853900 non-null    float64
dtypes: float64(2)
memory usage: 13.2 MB
```

```
In [82]: 1 df.describe()
```

```
Out[82]:
```

	sal	tem
count	817509.000000	853900.000000
mean	33.840350	10.799677
std	0.461843	4.243825
min	28.431000	1.440000
25%	33.488000	7.680000
50%	33.863000	10.060000
75%	34.196900	13.880000
max	37.034000	31.140000

```
In [83]: 1 df.fillna(method='ffill')
```

```
Out[83]:
```

	sal	tem
0	33.4400	10.500
1	33.4400	10.460
2	33.4370	10.460
3	33.4200	10.450
4	33.4210	10.450
...	...	...
864858	33.4083	18.744
864859	33.4083	18.744
864860	33.4150	18.692
864861	33.4062	18.161
864862	33.3880	17.533

864863 rows × 2 columns

```
In [84]: 1 df.fillna(value=0,inplace=True)
```

C:\Users\teppa\AppData\Local\Temp\ipykernel\_1224\1434098079.py:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))  
df.fillna(value=0,inplace=True)

```
In [85]: 1 df.isnull().sum()
```

```
Out[85]: sal    0  
tem    0  
dtype: int64
```

```
In [86]: 1 x=np.array(df['sal']).reshape(-1,1)  
2 y=np.array(df['tem']).reshape(-1,1)
```

```
In [87]: 1 df.dropna(inplace=True)
```


C:\Users\teppa\AppData\Local\Temp\ipykernel\_1224\1379821321.py:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))  
df.dropna(inplace=True)

```
In [88]: 1 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
2 reg=LinearRegression()
3 reg.fit(x_train,y_train)
4 print(reg.score(x_test,y_test))
```


0.014828355646333113

```
In [89]: 1 y_pred=reg.predict(x_test)
2 plt.scatter(x_test,y_test,color='b')
3 plt.plot(x_test,y_pred,color='k')
4 plt.show()
```



```
In [90]: 1 df500=df[:][:500]
2         sns.lmplot(x='sal',y='tem',data=df500,order=1,ci=None)
```

Out[90]: <seaborn.axisgrid.FacetGrid at 0x17c29516860>




```
In [91]: 1 from sklearn.linear_model import LinearRegression
2         from sklearn.metrics import r2_score
3         model=LinearRegression()
4         model.fit(x_train,y_train)
5         y_pred=model.predict(x_test)
6         r2=r2_score(y_test,y_pred)
7         print("R2 score:",r2)
```

R2 score: 0.014828355646333113

```
In [92]: 1 plt.figure(figsize = (10, 10))
        2 sns.heatmap(df.corr(), annot = True)
```

Out[92]: <Axes: >



```
In [93]: 1 features = df.columns[0:2]
        2 target = df.columns[-1]
        3 X = df[features].values
        4 y = df[target].values
        5 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=17)
        6 print("The dimension of X_train is {}".format(X_train.shape))
        7 print("The dimension of X_test is {}".format(X_test.shape))
        8 scaler = StandardScaler()
        9 X_train = scaler.fit_transform(X_train)
       10 X_test = scaler.transform(X_test)
```

The dimension of X\_train is (605404, 2)  
The dimension of X\_test is (259459, 2)



```
In [94]: 1 lr = LinearRegression()
2 lr.fit(X_train, y_train)
3 actual = y_test
4 train_score_lr = lr.score(X_train, y_train)
5 test_score_lr = lr.score(X_test, y_test)
6 print("\nLinear Regression Model:\n")
7 print("The train score for lr model is {}".format(train_score_lr))
8 print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 1.0

The test score for lr model is 1.0

```
In [95]: 1 ridgeReg = Ridge(alpha=10)
2 ridgeReg.fit(X_train,y_train)
3 train_score_ridge = ridgeReg.score(X_train, y_train)
4 test_score_ridge = ridgeReg.score(X_test, y_test)
5 print("\nRidge Model:\n")
6 print("The train score for ridge model is {}".format(train_score_ridge))
7 print("The test score for ridge model is {}".format(test_score_ridge))
```

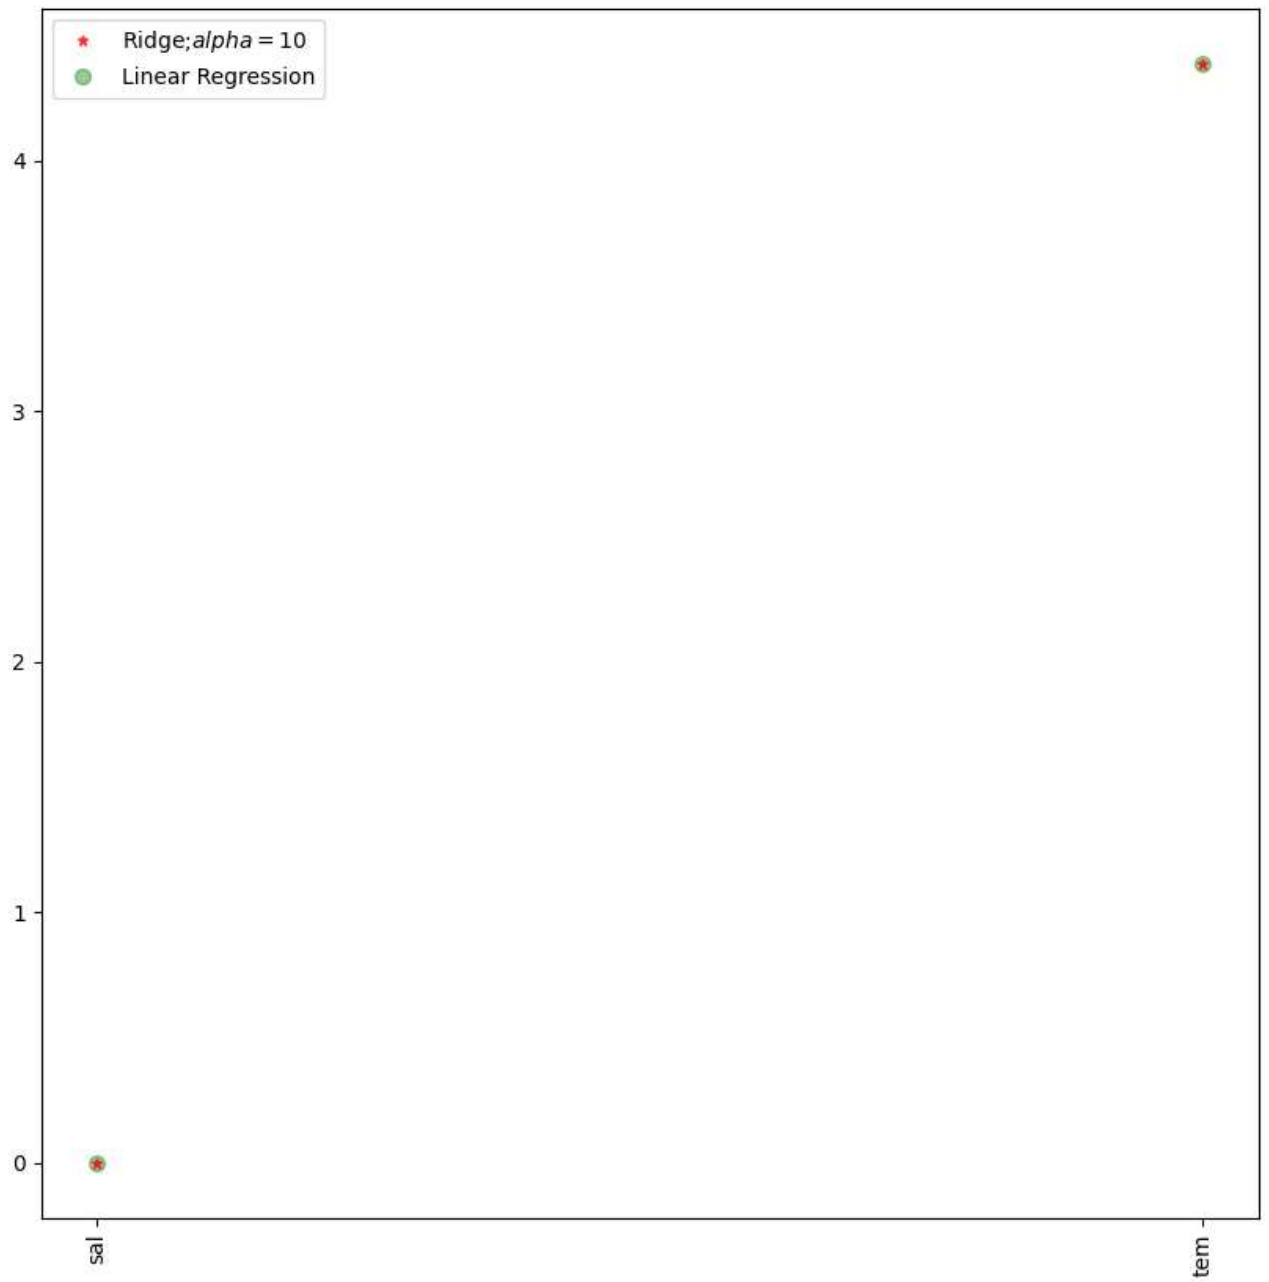
Ridge Model:

The train score for ridge model is 0.99999999723243

The test score for ridge model is 0.999999997231402

```
In [96]: 1 from sklearn.linear_model import Ridge,Lasso
2 from sklearn.preprocessing import StandardScaler
```

```
In [98]: 1 plt.figure(figsize=(10,10))
2 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='None',marker='*',markersize=5,color='red',label='Ridge;alpha = 10')
3 plt.plot(features,lr.coef_,alpha=0.4,linestyle="none",marker='o',markersize=7,color='green',label='Linear Regression')
4 plt.xticks(rotation=90)
5 plt.legend()
6 plt.show()
```



#Lasso Model

```
In [99]: 1 print("\nLasso Model: \n")
2 lasso = Lasso(alpha = 10)
3 lasso.fit(X_train,y_train)
4 train_score_ls =lasso.score(X_train,y_train)
5 test_score_ls =lasso.score(X_test,y_test)
6 print("The train score for ls model is {}".format(train_score_ls))
7 print("The test score for ls model is {}".format(test_score_ls))
```


Lasso Model:

The train score for ls model is 0.0

The test score for ls model is -1.9031696447013857e-05

```
In [100]: 1 pd.Series(lasso.coef_, features).sort_values(ascending = True).plot(kind = "bar")
```

Out[100]: <Axes: >



```
In [101]: 1
2 from sklearn.linear_model import LassoCV
3
4 lasso_cv = LassoCV(alphas = [0.0001, 0.001,0.01, 0.1, 1, 10], random_state=0).fit(X_train, y_train)
```

```
In [102]: 1
2 print(lasso_cv.score(X_train, y_train))
3 print(lasso_cv.score(X_test, y_test))
```

0.9999999994806811

0.9999999994806712

```
In [ ]: 1
```